



SMART CLASSROOM

A Mini Project Report

Submitted by

KARTHIKEYAN S

in association with AD23532 - Principles of Data Science

IN

ARTIFICIAL INTELLIGENCE AND DATA SCIENCE

RAJALAKSHMI ENGINEERING COLLEGE

[Autonomous]

Regulations 2023

RAJALAKSHMI ENGINEERING COLLEGE, CHENNAI 600 025

BONAFIDE CERTIFICATE

Certified that this report title “**SMART CLASSROOM**” is the Bonafide work of **KARTHIKEYAN S (2116231801081)** who carried out the mini project work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

SIGNATURE

Dr. J M GNANASEKAR,
HEAD OF THE DEPARTMENT,
Professor,
Department of AI&DS,
Rajalakshmi Engineering College
Chennai – 602 105.

SIGNATURE

SURENDAR A,
SUPERVISOR,
Assistant Professor,
Department of AI&DS,
Rajalakshmi Engineering College,
Chennai – 602 105.

Submitted for the Principles of Data Science Mini project review held on

Internal Examiner

External Examiner

ACKNOWLEDGEMENT

Initially I thank the Almighty for being with us through every walk of my life and showering his blessings through the endeavor to put forth this report.

My sincere thanks to our Chairman **Mr. S. MEGANATHAN, M.E., F.I.E.**, and our Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, M.E., Ph.D.**, for providing me with the requisite infrastructure and sincere endeavoring educating me in their premier institution.

My sincere thanks to **Dr. S.N. MURUGESAN, M.E., Ph.D.**, our beloved Principal for his kind support and facilities provided to complete our work in time.

I express my sincere thanks to **Dr. J M Gnanasekar M.E., Ph.D.**, Head of the Department of Artificial Intelligence and Data Science for his guidance and encouragement throughout the project work. I convey my sincere and deepest gratitude to our internal guide, **Surendar A**, Assistant Professor, Department of Artificial Intelligence and Data Science, Rajalakshmi Engineering College for his valuable guidance throughout the course of the project.

Finally, I express my gratitude to my parents and classmates for their moral support and valuable suggestions during the course of the project.

ABSTRACT

The **Smart Classroom Data Science Project** is an advanced educational technology system designed to enhance teaching and learning experiences through real-time analysis of student emotions and engagement. By integrating **data science, computer vision, deep learning and real-time analytics**, the system interprets facial expressions captured from classroom cameras to determine emotional states such as happiness, sadness, anger, fear, surprise, disgust, and neutrality. This information provides teachers with actionable insights to evaluate student attention levels and adapt instructional strategies accordingly.

Traditional classroom environments rely heavily on subjective assessments of engagement, often resulting in inaccurate evaluations of student participation. To address this limitation, the proposed system employs a **Convolutional Neural Network (CNN)** based on the **mini-XCEPTION architecture**, trained on the **FER2013 dataset**, which contains over 35,000 labeled grayscale facial images. The CNN uses **depthwise separable convolutions, batch normalization, and ReLU activations** to achieve real-time performance with efficient computation. The model is optimized using the **Adam optimizer** and **categorical cross-entropy loss function**, reaching approximately **66% validation accuracy** on FER2013 and maintaining reasonable performance in live testing scenarios.

The system's architecture consists of five main components: the input module for live video capture, preprocessing for face detection and normalization, the deep learning model for emotion classification, a backend server for processing and storage, and a **web-based dashboard** for visualization. Teachers can monitor emotion distributions, engagement ratios, and historical trends through an interactive interface that translates raw emotion data into meaningful engagement metrics.

This project demonstrates the potential of AI-driven emotion recognition in transforming traditional classrooms into **intelligent, responsive learning environments**. By providing continuous emotional feedback, the Smart Classroom system enables educators to identify disengaged students, adjust teaching methods, and promote active participation. Future enhancements include integrating multimodal data (voice tone, posture), applying reinforcement learning for adaptive recommendations, and implementing privacy-preserving techniques for secure data handling. Ultimately, the system highlights how **data science and deep learning-based emotion analytics** can foster more personalized, inclusive, and effective educational experiences.

Table of Contents

1. Abstract
2. Chapter 1: Introduction
 - 1.1 Problem Statement and Explanation
 - 1.2 Literature Survey
 - 1.3 Existing Systems
 - 1.4 Proposed System
3. Chapter 2: Data Collection and Preprocessing
 - 2.1 Data Collection and Description
 - 2.2 Workflow / Architecture Diagram
 - 2.3 Preprocessing Steps
4. Chapter 3: Results and Discussion
 - 3.1 Exploratory Data Analysis
 - 3.2 Algorithm Explanation
 - 3.3 Graphical Interface and Visualization
 - 3.4 Output
 - 3.5 Model Accuracy
 - 3.6 Conclusion
 - 3.7 Future Enhancements
5. References

CHAPTER -1: INTRODUCTION

1.1 Problem Statement and Explanation

In traditional classroom environments, the assessment of students' emotional and engagement levels has largely relied on the teacher's intuition and observation. While experienced educators can often perceive general patterns of attentiveness or disinterest, such assessments are inherently **subjective, inconsistent, and prone to bias**. Teachers may misinterpret subtle emotional cues, overlook quieter students, or fail to detect early signs of disengagement, especially in large or hybrid classrooms where continuous observation of every student is practically impossible.

Moreover, human perception alone cannot provide **quantitative or real-time data** about how students feel during lectures — whether they are happy, confused, frustrated, or bored. This lack of emotional awareness results in **delayed pedagogical interventions**, where teachers only realize disengagement after a test, feedback form, or declining performance. Consequently, teaching strategies often remain static and reactive rather than adaptive and responsive.

The absence of a **real-time feedback mechanism** between students and instructors reduces learning efficiency and motivation. Students who are emotionally disconnected may continue to struggle silently, while teachers remain unaware of the underlying causes. This challenge has become even more prominent with the rise of **digital and online learning**, where non-verbal cues are even harder to capture.

The **Smart Classroom Data Science Project** seeks to bridge this gap by introducing a **data science and deep learning–based emotion recognition and engagement analysis system**. By utilizing advanced computer vision and artificial intelligence techniques, the proposed system captures live video feeds from classroom cameras, detects student faces, and analyzes their facial expressions to infer emotional states in real time. The extracted emotional data is then processed to provide **actionable insights** into student engagement levels. Teachers can access these insights through a **web-based analytics dashboard**, enabling them to adjust their teaching strategies dynamically, ensure higher participation, and foster a more emotionally intelligent and inclusive learning environment.

1.2 Literature Survey

The field of **emotion recognition using artificial intelligence** has gained significant attention over the past decade, driven by breakthroughs in deep learning and computer vision. Traditional machine learning methods relied on handcrafted features such as edge detection, texture descriptors, and geometric measurements of facial landmarks. However, the introduction of **Convolutional Neural Networks (CNNs)** revolutionized image-based emotion analysis by allowing models to automatically learn hierarchical feature representations from raw image data.

Seminal work by **Goodfellow et al. (2016)** demonstrated the effectiveness of deep neural networks in complex visual recognition tasks. The emergence of datasets such as **FER2013**, **CK+**, **AffectNet**, and **JAFFE** provided large-scale labelled data for emotion recognition research. Studies utilizing these datasets have shown that CNNs, especially architectures like **VGGNet**, **ResNet**, and **Xception**, outperform traditional classifiers by significant margins.

Research by **Mollahosseini et al. (2017)** introduced *AffectNet*, which further validated the potential of deep learning for large-scale affective computing. Meanwhile, **Chollet (2017)** proposed the **Xception model**, based on depthwise separable convolutions, which improved computational efficiency while maintaining high accuracy — an approach particularly suitable for real-time applications such as classroom monitoring.

Despite these advancements, most emotion recognition systems remain **laboratory prototypes**, often tested under controlled lighting, fixed camera angles, and static subjects. Few systems address **real-time, real-world classroom conditions** characterized by variable lighting, multiple faces, occlusions, and dynamic environments. Furthermore, existing academic studies often focus purely on model performance rather than practical application and integration into educational ecosystems.

The **Smart Classroom Project** builds upon these foundational studies but distinguishes itself by focusing on **real-time implementation, system integration, and educational relevance**. It merges the domains of **affective computing, deep learning, and pedagogical analytics** to provide a holistic classroom monitoring solution.

1.3 Existing System

In current educational settings, classroom engagement monitoring typically relies on **manual observation**, periodic assessments, or student self-reports. Some schools and institutions employ video recording tools for classroom review, but these recordings serve primarily for administrative or training purposes rather than emotional analysis. Such methods are **reactive** rather than proactive — they only allow reflection after the session ends and provide no real-time insights to the instructor.

A few AI-based systems attempt to gauge sentiment through **text-based sentiment analysis** from feedback forms, chat responses, or digital participation logs. While this approach can reveal general opinions or satisfaction levels, it completely ignores the **non-verbal communication cues** — such as facial expressions, gaze direction, and micro-expressions — which are critical indicators of genuine emotion and engagement.

Moreover, these systems often lack the ability to correlate emotional data with **learning outcomes**, meaning teachers cannot easily interpret whether emotional fluctuations during lessons affect performance. The absence of integration between emotion analysis, student analytics, and teaching tools limits the potential of these systems in real educational environments.

Therefore, the **existing systems** fail to meet three essential requirements of an intelligent classroom:

1. **Real-time emotion detection** from visual cues.
2. **Quantitative engagement analytics** linked with learning activities.
3. **Actionable reporting tools** that help teachers adapt dynamically.

These limitations underscore the necessity for a **comprehensive, AI-driven classroom monitoring solution** that can function efficiently in real-world conditions.

1.4 Proposed System

The **proposed Smart Classroom System** introduces an innovative framework that integrates **real-time facial emotion recognition** with classroom engagement analytics. The system captures live video streams through cameras installed in the classroom. Using **deep learning–based computer vision techniques**, it identifies each student’s face, preprocesses the image (by converting it to grayscale and resizing it to 48×48 pixels), and classifies the emotion using a trained **mini-XCEPTION CNN model**.

The **mini-XCEPTION architecture** is a lightweight adaptation of the original Xception model. It employs **depth wise separable convolutions**, **batch normalization**, and **ReLU activations** to ensure high accuracy with reduced computational overhead — ideal for real-time operation on moderate hardware. The model is trained on the **FER2013 dataset**, which contains over 35,000 labelled images distributed across seven emotion classes: *Angry, Disgust, Fear, Happy, Sad, Surprise, and Neutral*.

The application follows a **client-server architecture**:

- The **client side** (camera module) captures and transmits the live video stream.
- The **server side** hosts the deep learning inference engine, processes emotion predictions, and stores them in a structured database.
- A **web-based dashboard** allows teachers to view aggregated analytics, such as emotion distribution charts, engagement scores, and historical trends.

This setup empowers educators to monitor not just attendance, but emotional participation. Teachers can identify patterns like prolonged sadness or confusion, indicating that certain topics may need re-teaching or a different instructional approach.

The system’s **real-time nature** ensures immediate feedback, allowing dynamic adjustment of lesson delivery. Additionally, data logging supports **long-term trend analysis**, helping institutions evaluate teaching effectiveness and student well-being over time.

In summary, the proposed Smart Classroom System represents a forward-looking approach that fuses **deep learning, affective computing, and educational data analytics** to create a **responsive, adaptive, and emotionally aware learning environment**.

2.1 Data Collection and Description

The foundation of any deep learning–based emotion recognition system lies in the **quality and diversity of the dataset** used for training and validation. For this project, the **Facial Expression Recognition 2013 (FER2013)** dataset has been chosen as the primary source of facial emotion data. The FER2013 dataset was originally introduced during the Kaggle Facial Expression Recognition Challenge and remains one of the most widely used benchmarks for training facial expression classifiers.

The dataset contains **35,887 labelled grayscale images** of human faces, each with a resolution of **48 × 48 pixels**. These images are categorized into **seven distinct emotion classes**, namely Angry, Disgust, Fear, Happy, Sad, Surprise, Neutral


Each image is carefully annotated to represent a specific emotional expression. The grayscale format ensures that the model focuses on **facial structure and muscle movement** rather than being influenced by colour variations or lighting conditions. This makes the dataset highly suitable for real-world applications where illumination and camera quality can vary significantly. The dataset is divided into **training, validation, and test subsets**, ensuring proper model generalization and preventing overfitting. During training, the system learns to identify discriminative features from facial components such as the **eyes, mouth, eyebrows, and cheeks**, which are critical indicators of human emotion.

In addition to this static dataset, the Smart Classroom system also collects **real-time video data** from classroom cameras during live sessions. This real-time data is used during the **inference phase**, where the trained model predicts the emotional states of students on live video frames. The combination of pre-trained models on FER2013 and live data testing ensures that the system functions effectively in both **controlled and dynamic classroom environments**.

Furthermore, real-time data contributes to continuous model improvement. By periodically retraining the model with captured classroom samples (subject to privacy and ethical guidelines), the system can adapt to the **specific demographics, cultural variations, and environmental conditions** of different classrooms. This hybrid approach—using both static datasets and real-time video streams—ensures that the Smart Classroom system remains accurate, flexible, and context-aware

2.2 Workflow Diagram (Architecture Diagram)

This **architecture diagram** illustrates a modular and scalable deep learning system for



classroom emotion analysis. Designed for real-time processing, it transforms video data into actionable insights for educators, with each module playing a distinct role in this efficient pipeline.

Input Module: Captures live video streams or processes stored footage from classroom cameras, segmenting video into frames for analysis.

Preprocessing Module: Prepares captured frames for the AI model by performing face detection, cropping, grayscale conversion, and normalization, ensuring clean and standardized input.

Deep Learning Model (mini-XCEPTION CNN): This core component receives pre-processed faces and classifies them into seven emotion categories, outputting probabilistic predictions of emotional states.

Backend Server: Aggregates emotion predictions, stores data in a database, and computes key engagement metrics and trends for comprehensive analysis.

Web Dashboard: Provides educators with an intuitive interface to visualize real-time emotion distributions, trend reports, and historical data, serving as a vital decision-support layer for classroom engagement.

2.3 Preprocessing steps

Face Detection: Locates and isolates facial regions within the video frame using tools like


Haar Cascade or Dlib. The detected face is cropped and uniformly resized to pixels to ensure consistent input for the model.

Grayscale Conversion: The cropped color image is converted to grayscale. This reduces the data's dimensionality and computational cost while preserving the essential structural features (contours, shadows) critical for emotion recognition.

Image Normalization: Pixel intensity values are standardized by dividing them by 255. This scales all values to a range between 0 and 1, which helps accelerate the neural network's convergence during training.

Data Augmentation: Applied to the training dataset to improve model generalization and prevent overfitting. Techniques like random horizontal flipping, rotation, zooming, and slight translations simulate real-world variations.

Label Encoding: Converts text-based emotion labels (e.g., "Happy") into categorical numeric formats using one-hot encoding. This creates binary vectors necessary for the model's output layer and for calculating categorical cross-entropy loss.



CHAPTER -3: RESULTS & DISCUSSIONS

3.1 Exploratory Data Analysis (EDA)

Before training the model, an **Exploratory Data Analysis (EDA)** was conducted on the **FER2013 dataset** to better understand its structure, distribution, and potential biases. The dataset contains **35,887 grayscale images** divided across seven emotion categories: *Angry*, *Disgust*, *Fear*, *Happy*, *Sad*, *Surprise*, and *Neutral*.

An analysis of emotion frequencies revealed that while most categories are fairly balanced, the **‘Disgust’** class contains significantly fewer samples compared to others. This class imbalance can influence the learning process, as the model might tend to favor classes with more examples (e.g., “Happy” or “Neutral”). To mitigate this issue, **data augmentation techniques** such as rotation, flipping, and zooming were applied, artificially increasing the number of samples for underrepresented emotions.

During visualization, emotion distributions were plotted using bar charts and confusion matrices to identify potential overlaps between similar emotions. For instance, “Sad” and “Fear” expressions often share subtle similarities in eye and mouth movement patterns, leading to occasional misclassifications.

The **facial features** that play the most significant role in emotion classification include:

- **Eye shape and openness** (e.g., wide-open eyes in “Surprise” vs. narrowed eyes in “Anger”)
- **Mouth curvature** (upward in “Happy”, downward in “Sad”)
- **Eyebrow position and tilt**, which help distinguish between “Anger” and “Disgust”.

These insights confirm that the dataset provides a rich set of discriminative features for the model to learn meaningful emotion patterns.

3.2 Algorithm Explanation

The Smart Classroom project employs a **deep learning–based Convolutional Neural Network (CNN)** known as **mini-XCEPTION**, a lightweight and efficient variant of the Xception architecture proposed by **François Chollet (2017)**. This architecture is designed for **real-time emotion recognition** with reduced computational cost while maintaining high accuracy.

Model Architecture Overview

The mini-XCEPTION model leverages:

- **Depthwise Separable Convolutions:** These reduce the number of parameters and computations by separating spatial and channel-wise feature extraction.
- **Batch Normalization:** Stabilizes and accelerates training by normalizing feature maps.
- **ReLU Activation:** Introduces non-linearity, allowing the model to learn complex patterns.
- **Global Average Pooling:** Replaces fully connected layers, reducing overfitting and improving generalization.
- **Softmax Output Layer:** Provides probability distributions across seven emotion classes.

The model is trained using **categorical cross-entropy loss** and optimized using the **Adam optimizer**, which adapts learning rates during training. **Transfer learning** is also utilized by initializing weights from a pre-trained model, improving convergence and accuracy.

Model Training Process

1. Input images are resized to **48×48 pixels** and normalized.
2. The model processes each image through a series of convolutional and separable convolutional layers.
3. The learned features are pooled and passed through a dense softmax classifier.
4. The model minimizes loss through backpropagation and updates weights using Adam optimization.

Sample Code for the mini-XCEPTION Model

Below is a simplified code snippet demonstrating how the mini-XCEPTION CNN model was built and trained using **TensorFlow/Keras**:

```
from tensorflow.keras.models

import Sequential

from tensorflow.keras.layers import Conv2D, SeparableConv2D, BatchNormalization,
ReLU, GlobalAveragePooling2D, Dense

num_classes = 7

model = Sequential([

    Conv2D(32, (3,3), padding='same', input_shape=(48,48,1)),

    BatchNormalization(),

    ReLU(),

    SeparableConv2D(64, (3,3), padding='same'),

    BatchNormalization(),

    ReLU(),

    SeparableConv2D(128, (3,3), padding='same'),

    BatchNormalization(),

    ReLU(),

    GlobalAveragePooling2D(),

    Dense(num_classes, activation='softmax')

])

model.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])

model.summary()
```


Once trained, the model is saved in **HDF5 format (.h5)** for deployment:

```
model.save("mini_xception_emotion_model.h5")
```

This trained model is later integrated into the Smart Classroom backend for **live emotion inference**.

3.3 Graphical Diagram (Interface and Visualization)

The **graphical user interface (GUI)** of the Smart Classroom System provides a real-time visualization of emotion predictions and engagement statistics. The dashboard is built using **web-based technologies (HTML, Flask, Chart.js, and Python backend)**, offering educators a simple and interactive experience.

Key graphical components include:

- **Live Emotion Bar Chart:** Displays the percentage of detected emotions (e.g., 40% Happy, 30% Neutral, 10% Sad, etc.) for the current session.
- **Engagement Analytics Panel:** Summarizes metrics like *Attention Ratio*, *Emotional Stability Index*, and *Dominant Emotion Trend*.
- **Historical Reports:** Track emotion distribution trends across multiple lectures to analyze engagement patterns over time.

These visualizations allow teachers to instantly interpret student behavior and adapt teaching approaches dynamically — for instance, revisiting difficult concepts when “Confused” or “Sad” expressions dominate.



3.4 Output

Face Confused

Class Insights (Inbuilt Analysis)

Analyze Now Class: Great

```
{
  "overall": "great",
  "summary": "Student appears focused and neutral",
  "advice": "Continue monitoring student behavior and provide support as needed"
}
```

Teacher Alerts & Actions

Attention: okay Alert: Active

Student: Student shows signs of frustration
Action: Check if student needs help, consider breaking down the task into smaller steps
Emotion: angry (51%) - 0s ago

```
{
  "stress_level": "0.30",
  "fatigue_level": "0.50",
  "suggestion": "You're maintaining a healthy pace.",
  "student_state": "distracted",
  "confidence": "37%"
}
```

Primary Student Status

LOCKED STUDENT (ID: 1)
Student shows signs of frustration
Emotion: angry (51%)
Status: frustrated (high severity)
Session Active - Tracking Only This Student

Recent Detections


```
[
  {
    "t": 1760505969214,
    "faces": [
      {
        "id": 1,
        "emotion": "surprise",
        "confidence": 0.3932291567325592,
        "is_confused": false
      }
    ]
  },
  {
    "t": 1760505969214,
    "faces": [
      {
        "id": 1,
        "emotion": "surprise",
        "confidence": 0.3932291567325592,
        "is_confused": false
      }
    ]
  }
]
```

Realtime Emotion Detection

API: http://localhost:8000

Start Camera Stop Camera Start Realtime Stop Realtime

PRIMARY STUDENT - surprise 30%



Backend: OK Camera: active Realtime: active FPS: 2 Faces: 1

3.5 Model Accuracy

The model achieved an average **validation accuracy of approximately 66%** on the FER2013 dataset. This performance aligns with reported results in similar deep learning emotion recognition studies. The confusion matrix revealed high precision for “Happy” and “Neutral” emotions, while “Fear” and “Disgust” showed moderate overlap due to visual similarity.

In real-world classroom scenarios, accuracy slightly fluctuates depending on factors such as:

- **Lighting conditions and camera resolution**
- **Student face angles and occlusions (e.g., masks or glasses)**
- **Number of faces in frame (multi-face detection load)**

Despite these challenges, the system maintained reliable emotion detection performance suitable for live monitoring. Ongoing fine-tuning, adaptive preprocessing, and continuous retraining on classroom-collected data are expected to improve robustness further.

3.6 Conclusion

The Smart Classroom Data Science Project demonstrates the successful application of **AI-driven facial emotion recognition** to enhance educational environments. By integrating deep learning with classroom analytics, the system enables real-time monitoring of student engagement, helping teachers make data-informed pedagogical decisions.

The project validates that **CNN-based emotion recognition models** like mini-XCEPTION can effectively capture subtle facial cues and translate them into measurable engagement indicators. Furthermore, by combining model predictions with web-based analytics, the system transforms abstract emotion data into **practical insights** that directly improve teaching efficiency and learner participation.

Overall, the results highlight the transformative potential of **AI-powered Smart Classrooms** in fostering responsive, adaptive, and emotionally intelligent education systems.

3.7 Future Enhancement

Although the current implementation achieves promising results, several enhancements can extend its capabilities and performance in future versions:

1. **Expanded Emotion Categories:** Include complex affective states like *boredom*, *confusion*, *curiosity*, and *frustration* to capture richer emotional dynamics.
2. **Multimodal Emotion Analysis:** Combine visual data with **voice tone**, **posture**, and **eye-tracking** for holistic emotion understanding.
3. **Cloud and Mobile Integration:** Develop cloud-based and mobile-compatible applications to support remote learning environments.
4. **Reinforcement Learning:** Implement adaptive feedback mechanisms to automatically recommend teaching interventions based on engagement levels.
5. **Privacy and Ethics:** Strengthen data anonymization, encryption, and ethical safeguards to ensure privacy in emotion-based analytics.

References

- [1] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*, MIT Press, 2016.
- [2] A. Mollahosseini, B. Hasani, and M. H. Mahoor, "AffectNet: A Database for Facial Expression, Valence, and Arousal Computing in the Wild," *IEEE Transactions on Affective Computing*, vol. 10, no. 1, pp. 18–31, 2017.
- [3] S. Li and W. Deng, "Deep Facial Expression Recognition: A Survey," *IEEE Transactions on Affective Computing*, 2020, pp. 1–1.
- [4] FER2013 Dataset – Kaggle Facial Expression Recognition Challenge. [Online]. Available: <https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge/data>
- [5] F. Chollet, "Xception: Deep Learning with Depthwise Separable Convolutions," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 1800–1807.
- [6] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, et al., "Generative Adversarial Nets," in *Advances in Neural Information Processing Systems*, vol. 27, 2014, pp. 2672–2680.
- [7] Dlib C++ Library: Machine Learning Toolkit and Face Detection API. [Online]. Available: <http://dlib.net/>